


Adversarial Search



These slides are based on the slides created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley - <http://ai.berkeley.edu>.

The artwork is by Ketrina Yim.

Academic Integrity

It is NOT OK:

- ✗ to share code, or copy code from fellow students or from the web.
- Infractions will be detected and will lead to an automatic 0 on the given assignment or test.
- Infractions will also be reported to the Office of Student Conduct and Ethical Development.

Academic Integrity

You copy and paste code from the web or from another student

⇒ Automatic 0 on the assignment and report to the Office of Student Conduct and Ethical Development.

You copy code from the web or from another student but you change some variable names.

⇒ Automatic 0 on the assignment and report to the Office of Student Conduct and Ethical Development.

Academic Integrity


You see someone else's code (from a friend or on the web)

To avoid an academic integrity violation, you must:

1. **Cite the source** that you saw - web address or friend's name
2. Rewrite the code with your own 'words' **without looking at the source.**
3. **Comment** extensively so that you explain what every line of code does.

Game Playing State-of-the-Art

- **Checkers:** 1950: First computer player.
1994: First computer champion: Chinook
used 8 piece endgame database
2007: Checkers solved!
- **Chess:** 1997: Deep Blue defeated human
champion Gary Kasparov. Current
programs are even better, if less historic.
- **Go:** $b > 300$. Oct 2015: Google DeepMind
program AlphaGo defeated Fan Hui, the
European Go champion. It uses a
combination of machine learning and
randomized tree search techniques.
- **Pacman**



Adversarial Games



Types of Games

- Many different kinds of games!
- Axes:
 - Deterministic or stochastic?
 - One, two, or more players?
 - Zero sum?
 - Fully observable?



Games


How are games different from the search problems we've seen so far?

- so far we have been looking for a sequence of actions, that are guaranteed to succeed
- that does not work here because we don't control our opponent
- we need to be able to decide on a move from each state


We need algorithms for calculating a **strategy (policy)** which recommends a move from each state

Games

- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P=\{1...N\}$ (usually take turns)
 - Actions: A (may depend on player / state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test (game over?): $S \rightarrow \{t, f\}$
 - Terminal Utilities: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$



Zero-Sum Games



Zero-Sum Games


- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition

General Games


- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible

Adversarial Search

- We still need to think about the consequences of our actions.
- We also need to think about our opponent and how they will respond to our actions.




Single-Agent Trees



Value of a State

Value of a state:
The best achievable
outcome (utility)
from that state




Non-Terminal States:

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:

$$V(s) = \text{known}$$

Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree



MAX (X)



MIN (O)




MAX (X)



MIN (O)

TERMINAL

Utility



minimax value
at the root?


A. -1

B. 0

C. 0.5

D. 1

E. 2




Adversarial Search (Minimax)

Deterministic, zero-sum games:

- One player maximizes result
- The other minimizes result

Minimax search:

- A state-space search tree
- Players alternate turns
- Compute each node's **minimax value**: the best achievable utility against a rational adversary



Terminal values:
part of the game

Minimax Implementation

def value(state):

if the state is a terminal state: return the state's utility

if the agent is MAX: return max-value(state)

if the agent is MIN: return min-value(state)

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

def min-value(state):



initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return v




Minimax Example



Minimax Properties

- Complete?
 - Yes, if tree is finite
- Optimal?

Minimax Properties



Optimal against a perfect player. Otherwise?